


Design quality that keeps pace with how fast engineering ships.

At realfast, most user-facing interface code is now written by AI agents. I built exo-designer to give those agents the design system before they write a line, so every screen is on-brand and accessible the first time, with no review step and no designer added to the pod.



Production software in weeks, with senior engineers and frontier models.

An AI-native software delivery firm: legacy modernisation, AI agent workflows, Salesforce implementations, and AI transformation engagements, on an in-house platform called Exo. Pods of two to five senior engineers, agents doing most of the building, outcomes tied to metrics agreed before work starts. Clients include Z47, AVPN, TASC, Houzbay, BoardRoom, and Hypernative.

~60%

gross margin, as reported

26×

revenue growth in 2026, as reported

BACKED BY



The founders built and sold one engineering firm, then ran product and technology at Gojek.



Sidu Ponnappa

CEO, Co-founder

Co-founded C42 Engineering, acquired by Gojek in 2015. Then CIO of Gojek Singapore, MD of Gojek India, and a Gojek Group board member.



Aakash Dharmadhikari

CPTO, Co-founder

Co-founded C42. Group Chief Product Officer and board member at Gojek. Later CPO at Jiva.




Steve Sule

CFO, Co-founder

Former Partner at Edith Grove Capital, previously Arma Partners and M&A at Mahindra & Mahindra. Joined as an advisor, became a co-founder.

One designer, many engineering pods, and every pod shipping a slightly different brand.



No designer touches most of the screens realfast ships.

My job was to build design production infrastructure and pipelines for an engineering organisation, so that every one of those screens comes out well designed anyway.

Side by side, you could see the drift.

When I joined, every internal tool and client interface was built by whoever needed it, usually with an AI assistant writing the code. Each one worked fine on its own, but colours, form controls, and spacing changed from screen to screen. Nobody was being careless. There was simply no standard for the tools to follow.

I wanted LLMs to produce realfast's brand, at the same speed.

They can produce a settings page in under a minute. Without guidance, they produce the average of every settings page they have ever seen, which is nobody's brand.

WHAT THE ORGANISATION GAINS

01

A brand that holds everywhere

Internal tools end up in client demos, screenshots, and hiring loops. It matters that they all look like the same company built them.

02

Accessibility without a checklist

Labels, named controls, and contrast are WCAG basics, and a real liability on client work. When they are the default, nobody has to remember them.

03

Review time back to the team

By the time I review a finished tool, I am redesigning it. Skipping that step saves the pod a day and frees me for work only a designer can do.

Instead of checking screens after they were built, I gave the AI the rules before it started.

AT FIRST

Request

AI writes UI

Designer reviews

Rework

Ship, late

↑

each review sent the work back here

My first attempt was a linter, and it changed nothing upstream.

It scanned, flagged, and fixed what it could. Each audit found the same issues, we cleaned them up, and the next feature brought them back.

NOW

CLAUDE.md is read at the start of every session

Rules in project

Request

AI writes governed UI

Self-check

Ship

Now the first draft is the one that follows the rules.

I moved the design system into the file the model reads first: the rules, examples of right and wrong, and a checklist it runs before it shows any code. The audit is still there, but as a check rather than a gate.

- WHAT I HELD TO
- Put the rules where the code is written

Show examples, not just rules

Check on request, not on every screen

I directed the build, Claude Code wrote the code, and I cut what we did not need.

Rules as a project file, not a Figma library

Engineers and agents never open Figma. The system had to live where code is written.

Built a UI compliance dashboard, deleted it three days later

Nobody opened it. If the report is not in the terminal, nobody reads it.

Graded the audit on usability, not lint rules

Six weighted categories, feedback states counting most. A screen with no loading state is broken even when every colour is right.

One set of skills for React, Salesforce, and Rails

The tool detects the stack. Engineers learn one set of skills and use it on every project.

I DECIDED

- 01 **The thesis:** rules before the code, not review after it
- 02 **What counts as a violation,** and how much each category weighs
- 03 **The role of every brand colour,** all sixteen of them
- 04 **Killing the dashboard** and staying in the terminal
- 05 **The wording of every rule,** since the model does exactly what the text says

CLAUDE CODE PRODUCED

- 01 **Scanners** for accessibility, theming, and component usage
- 02 **Auto-fixers** for colours, spacing, and labels
- 03 **Seventeen pattern templates,** to my specification
- 04 **The CLI, installer,** and generator plumbing
- 05 **First-draft documentation,** which I rewrote

Eight skills that do three jobs. Engineers work the way they always did, and the screens come out right.

- 1

Set up

ONCE PER PROJECT

Detects the framework, writes the theme and CSS variables, drops the design rules into the project.

/design-init /design-theme
- 2

Make

EVERY DAY

New and rebuilt screens follow the rules, built from 17 reference patterns.

/design-generate /design-redesign
- 3

Check

WHEN ASKED

A second opinion when someone wants one: a quick pass, a weighted audit, an auto-fix, and a team report.

/design-check /design-audit /design-fix /design-report

```
claude · internal-tool

$ npm install exo-designer
  exo-designer: skills installed. Next: type /design-init

> /design-init
  Detected React + Tailwind. Applied preset: RealFast Branding.
  Wrote .design-agent/theme.json, src/styles/theme.css, CLAUDE.md

> Build a user management page with a table, search, and filters.
  Using patterns: page-layout, data-table, search-filter, empty-state
  Self-check  components ✓  tokens ✓  labels ✓  focus ✓  loading ✓
              empty ✓  contrast ✓  spacing ✓  errors ✓  responsive ✓
  Created src/pages/users.tsx

> /design-audit
  0 violations across 14 files. Score 100/100.
```

example session · commands and output as they appear in Claude Code

O

hardcoded colours, native inputs, or unlabelled fields on any screen. The rules stop them before anyone reviews.

NOT MEASURED YET

How many review hours it saved, and how many teams use it. The numbers above describe the system, not its effect. Next I want to audit tools that shipped before and after it, add checks for error recovery, and run it on client work with the client's own palette.

What it delivers today, what I have not measured yet, and where it goes next.

-
- | | |
|----|---|
| 6 | quality checks on every screen
weighted toward what users notice first: feedback, forms, brand |
| 17 | reference patterns
for the screens engineers build most often |
| 3 | frameworks from one set of skills
React, Salesforce, and Rails, with the same skills |
-

